

Full NR and LASSO Run

Tucker Morgan - tlm2152

3/26/2022

This file prepares our data, runs Newton-Raphson, runs LASSO, and plots some results at the bottom. I think this file should stay clean and relatively simple. I think the file should only take 15ish minutes to knit or run in full. I think the best idea is for plots to be saved as R objects and imported into the presentation / report, but open to ideas here.

```
source("./shared_code/setup.R")
knitr::opts_chunk$set(cache = TRUE)
```

```
source("./shared_code/data_prep.R")
```

Newton-Raphson

```
source("./shared_code/full_NR2.R")
source("./shared_code/roc_func.R")
roc_nr <- roc_func(nr_beta_est, bc_tst)
```

LASSO

```
source("./shared_code/logistic_lasso.R")
source("./shared_code/auc_calc_lasso.R")
```

Folding Data

```
source("./shared_code/cv_folding.R")
```

```
bc_trn_folds <-
  cv_sets(training = bc_trn) %>%
  select(-fold_p)
```

```
source("./shared_code/cv_implementation.R")
lambda_max <- max(t(scale(as.matrix(bc_trn[, -1])))) %*% bc_trn[, 1] / nrow(bc_trn[, -1])
lambda_list <- list(log(lambda_max), log(0.0001), -(log(lambda_max) - log(0.0001))/100, exp)
```

```
cv_res <- cv_jt(k = 5, training = bc_trn_folds,
  func = logistic_lasso, lam_start_stop_func = log,
  lambda_list = lambda_list)
```

```
# pulling out key results
selected_lambda <-
  cv_res[[1]][[length(cv_res[[1]])]] %>%
```

```

  filter(mean_auc == max(mean_auc)) %>%
  pull(lambda) %>%
  mean()

lasso_final_model <- logistic_lasso(inputs = bc_trn[,-1],
                                   output = bc_trn[,1],
                                   lambda_vec = selected_lambda)
lasso_betas <- lasso_final_model[[2]] %>% t()
# this is just for viz purposes
tst_lambda_vec <- exp(seq(from = log(lambda_max),
                        to = log(0.0001),
                        by = -(log(lambda_max) - log(0.0001))/100))

lasso_final_range <- logistic_lasso(inputs = bc_trn[,-1],
                                   output = bc_trn[,1],
                                   lambda_vec = tst_lambda_vec)

lfr_df <- data.frame(do.call(cbind, lasso_final_range)) %>%
  select(-selected) %>%
  pivot_longer(cols = starts_with("beta"),
              names_prefix = "beta.",
              names_to = "beta_coef",
              values_to = "coef_est")


# storing lasso ROC object
roc_lasso <- roc_func(lasso_betas, bc_tst)

cv_res_lam <- cv_res[[1]]

# each of these plots could be saved as R objects and imported into other documents
lfr_df %>%
  group_by(lambda) %>%
  filter(beta_coef != "intercept") %>%
  ggplot(x = lambda, y = coef_est, group = beta_coef) +
  geom_path(aes(x = lambda, y = coef_est, group = beta_coef, col = beta_coef)) +
  coord_cartesian(xlim = c(0, 0.05)) + # this zooms in on the plot, comment out if flipping
# scale_x_reverse() + # this flips the x-axis if wanted
# coord_cartesian(xlim = c(0.06, 0)) + # this zooms in if flipped
  geom_vline(xintercept = selected_lambda) +
  labs(title = "Beta Coefficients in LASSO Model")

```

Beta Coefficients in LASSO Model

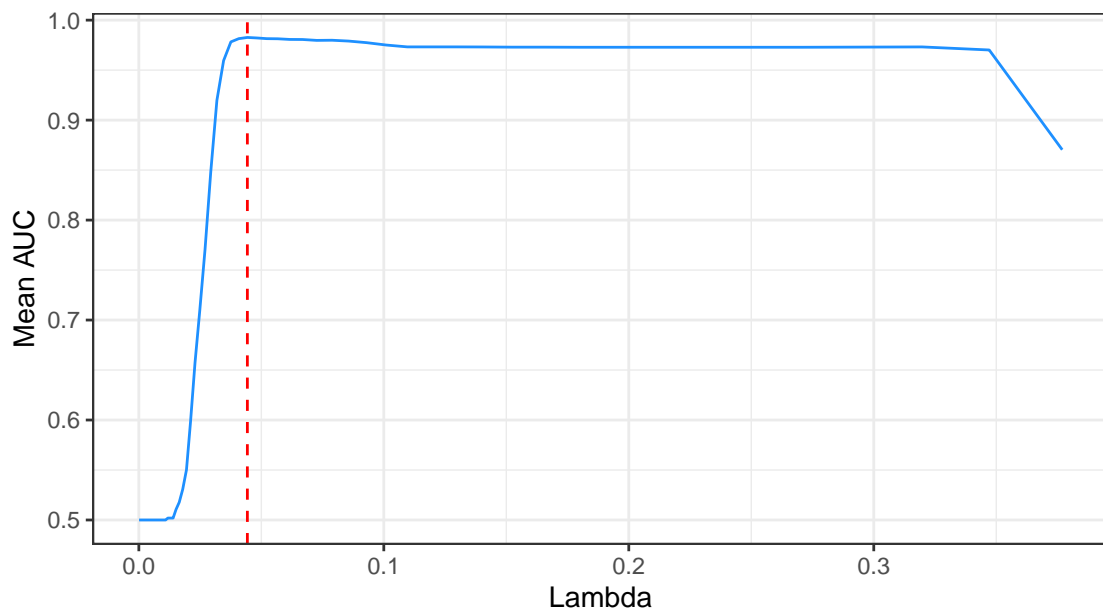


compactness_mean	concave.points_se	fractal_dimension_se	smoothness_me
compactness_se	concavity_se	fractal_dimension_worst	smoothness_se
compactness_worst	concavity_worst	radius_se	smoothness_wo
concave.points_mean	fractal_dimension_mean	radius_worst	symmetry_mean

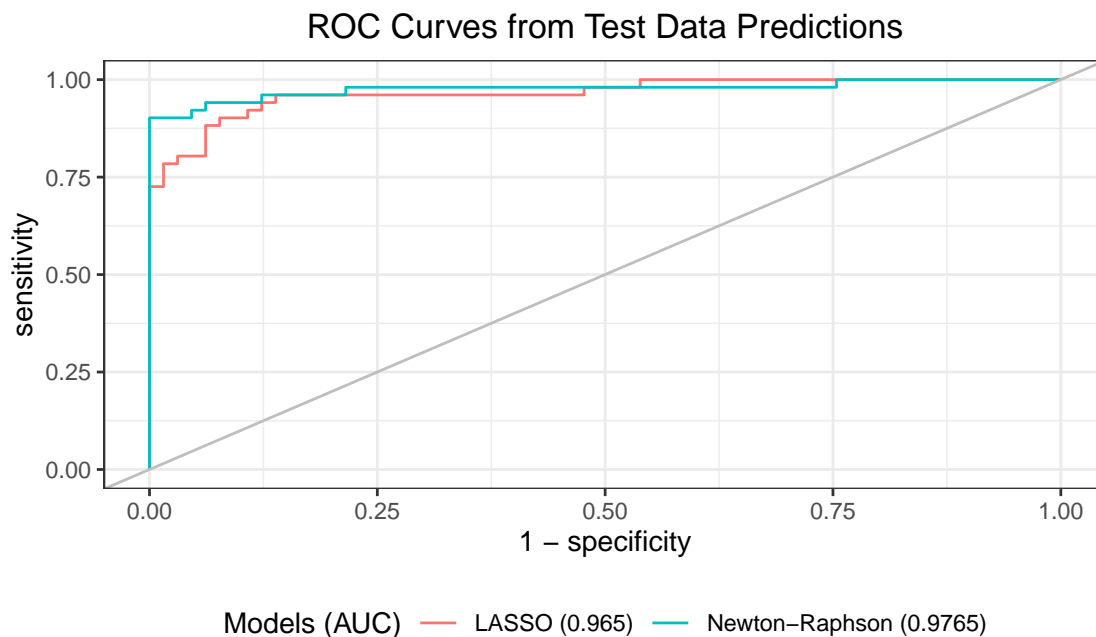
this looks alright, might prefer a larger lambda range to show AUC down to 0.5
but probably not important

```
cv_res[[1]][[1]] %>%
  data.frame() %>%
  ggplot(x = lambda, y = mean_auc) +
  geom_line(aes(x = lambda, y = mean_auc), col = "dodgerblue") +
  geom_vline(xintercept = selected_lambda, linetype = "dashed", col = "red") +
  labs(title = "Mean AUC vs. Lambda",
       x = "Lambda",
       y = "Mean AUC")
```

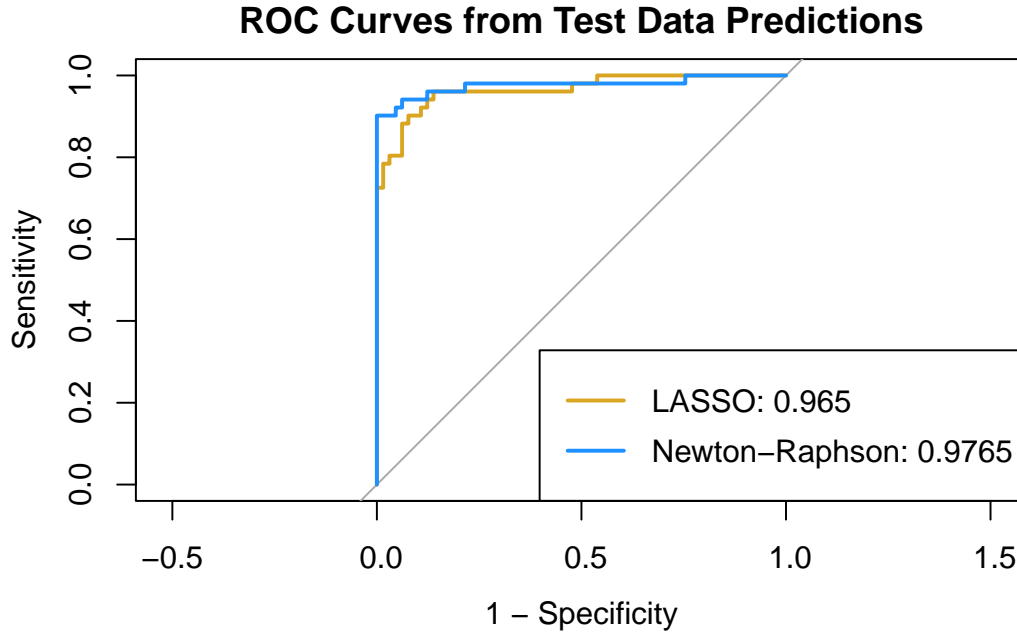
Mean AUC vs. Lambda



```
# Final ROC plots
auc_vec <- c(round(roc_lasso$auc[1], digits = 4), round(roc_nr$auc[1], digits = 4))
model_names <- c("LASSO", "Newton-Raphson")
ggroc(list(roc_lasso, roc_nr), legacy.axes = TRUE) +
  scale_color_discrete(labels = paste0(model_names, " (", auc_vec, ")"),
    name = "Models (AUC)") +
  geom_abline(intercept = 0, slope = 1, color = "grey") +
  labs(title = "ROC Curves from Test Data Predictions")
```



```
## another option for plotting ROC curves
plot(roc_lasso, legacy.axes = TRUE, col = "goldenrod",
  main = "ROC Curves from Test Data Predictions")
plot(roc_nr, col = "dodgerblue", add = TRUE)
legend("bottomright", legend = paste0(model_names, ": ", auc_vec),
  col = c("goldenrod", "dodgerblue"), lwd = 2)
```



```
beta_names <-
  (dimnames(lasso_betas)[1]) %>%
  data.frame() %>%
  rename(beta_coef = 1)

cbind(nr_beta_est, lasso_betas) %>%
  data.frame() %>%
  round(digits = 4) %>%
  rename(NewtonRaphson = 1, LASSO = 2) %>%
  knitr::kable(caption = "Final Beta Coefficient Estimates")
```

Table 1: Final Beta Coefficient Estimates

	NewtonRaphson	LASSO
intercept	-162.1162	-0.7872
texture_mean	1.9226	0.3641
smoothness_mean	-180.9337	0.0000
compactness_mean	-206.9520	0.0000
concave points_mean	539.5020	0.9912
symmetry_mean	-71.9190	0.0000
fractal_dimension_mean	555.5034	0.0000
radius_se	66.8227	0.0000
texture_se	-3.7168	0.0000
smoothness_se	424.7338	0.0000
compactness_se	264.4495	0.0000
concavity_se	-215.0474	0.0000
concave points_se	1443.8850	0.0000
symmetry_se	-1672.7126	0.0000
fractal_dimension_se	-4616.3556	0.0000
radius_worst	1.9170	1.4934
smoothness_worst	29.3660	0.0702
compactness_worst	-91.4676	0.0000

	NewtonRaphson	LASSO
concavity_worst	71.7133	0.2184
symmetry_worst	198.9776	0.1473
fractal_dimension_worst	429.9555	0.0000