

Untitled

Amy Pitts

3/17/2022

Data import

This dataset has 569 rows and 32 columns. The outcome variable of interest is **Diagnosis** which takes on values benign or malignant cases. There are 357 benign cases and 212 malignant cases as seen below in Table 1. One variable is **id** and the rest 30 variables are the mean, sd and largest values of the following criteria.

- radius (mean of distances from center to points on the perimeter)
- texture (standard deviation of gray-scale values)
- perimeter
- area
- smoothness (local variation in radius lengths)
- compactness ($\text{perimeter}^2 / \text{area} - 1.0$)
- concavity (severity of concave portions of the contour)
- concave points (number of concave portions of the contour)
- symmetry
- fractal dimension ("coastline approximation" - 1)

Using this dataset models will be compared and compared on their ability to predict cancer diagnosis.

Characteristic	**Benign**, N = 357	**Malignant**, N = 212	**p-value**
radius_mean	12.147 (1.781)	17.463 (3.204)	<0.001
texture_mean	17.915 (3.995)	21.605 (3.779)	<0.001
perimeter_mean	78.075 (11.807)	115.365 (21.855)	<0.001
area_mean	462.790 (134.287)	978.376 (367.938)	<0.001
smoothness_mean	0.092 (0.013)	0.103 (0.013)	<0.001
compactness_mean	0.080 (0.034)	0.145 (0.054)	<0.001
concavity_mean	0.046 (0.043)	0.161 (0.075)	<0.001
concave.points_mean	0.026 (0.016)	0.088 (0.034)	<0.001
symmetry_mean	0.174 (0.025)	0.193 (0.028)	<0.001
fractal_dimension_mean	0.063 (0.007)	0.063 (0.008)	0.5
radius_se	0.284 (0.113)	0.609 (0.345)	<0.001
texture_se	1.220 (0.589)	1.211 (0.483)	0.6
perimeter_se	2.000 (0.771)	4.324 (2.569)	<0.001
area_se	21.135 (8.843)	72.672 (61.355)	<0.001
smoothness_se	0.007 (0.003)	0.007 (0.003)	0.2
compactness_se	0.021 (0.016)	0.032 (0.018)	<0.001
concavity_se	0.026 (0.033)	0.042 (0.022)	<0.001
concave.points_se	0.010 (0.006)	0.015 (0.006)	<0.001
symmetry_se	0.021 (0.007)	0.020 (0.010)	0.028
fractal_dimension_se	0.004 (0.003)	0.004 (0.002)	<0.001
radius_worst	13.380 (1.981)	21.135 (4.284)	<0.001
texture_worst	23.515 (5.494)	29.318 (5.435)	<0.001
perimeter_worst	87.006 (13.527)	141.370 (29.457)	<0.001
area_worst	558.899 (163.601)	1,422.286 (597.968)	<0.001
smoothness_worst	0.125 (0.020)	0.145 (0.022)	<0.001
compactness_worst	0.183 (0.092)	0.375 (0.170)	<0.001
concavity_worst	0.166 (0.140)	0.451 (0.182)	<0.001
concave.points_worst	0.074 (0.036)	0.182 (0.046)	<0.001
symmetry_worst	0.270 (0.042)	0.323 (0.075)	<0.001
fractal_dimension_worst	0.079 (0.014)	0.092 (0.022)	<0.001

Build a logistic model to classify the images into malignant/benign

Before building the model we need to first write down the likelihood function, its gradient and Hessian matrix.

The likelihood function for our data which has a single binary response and 30 numerical explanatory variables is

$$\pi_i = P(Y_i = 1 | x_{i,1}, \dots, x_{i,30}) = \frac{e^{\beta_0 + \beta_1 x_{i,1} + \dots + \beta_{30} x_{i,30}}}{1 + e^{\beta_0 + \beta_1 x_{i,1} + \dots + \beta_{30} x_{i,30}}} = \frac{e^{\beta_0 + \sum_{j=1}^{30} \beta_j x_{i,j}}}{1 + e^{\beta_0 + \sum_{j=1}^{30} \beta_j x_{i,j}}}$$

Where \mathbf{X}_i represents the i observation of all 30 of our predictor variables. For the data give we have the likelihood is given by

$$L(\mathbf{X}|\beta) = \prod_{i=1}^n [\pi_i^{y_i} (1 - \pi_i)^{1-y_i}]$$

Finding the log-likelihood we have

$$l(\mathbf{X}|\vec{\beta}) = \sum_{i=1}^n \left[y_i \left(\beta_0 + \sum_{j=1}^{30} \beta_j x_{i,j} \right) - \log \left(1 + \exp \left(\beta_0 + \sum_{j=1}^{30} \beta_j x_{i,j} \right) \right) \right]$$

The gradient can then can be solved for. Observe

$$\nabla l(\mathbf{X}|\vec{\beta}) = \begin{bmatrix} \sum_{i=1}^n y_i - \pi_i & \sum_{i=1}^n x_{i,1}(y_i - \pi_i) & \dots & \sum_{i=1}^n x_{i,30}(y_i - \pi_i) \end{bmatrix}^T_{(1 \times 31)}$$

Finally, with the gradient we can derive our hessian. Note that due to the 30 predictor variables the hessian will be a 31 by 31 matrix.

$$\begin{aligned} \nabla^2 l(\mathbf{X}|\vec{\beta}) &= - \sum_{i=1}^n \begin{pmatrix} 1 \\ X \end{pmatrix} (1 - X) \pi_i (1 - \pi_i) \\ &= - (1 \quad X) \text{diag}(\pi_i (1 - \pi_i)) \begin{pmatrix} 1 \\ X \end{pmatrix} \end{aligned}$$

Where $X = (x_{i,1}, \dots, x_{i,30})$. Note that this matrix will always be negative definite at all parameters making this a well behaved problem.

Develop a Newton-Raphson algorithm to estimate your model

Modifications:

- I include half stepping in the Newton-Raphson method.
- Assent direction

Data Preprocess

```
set.seed(7777)
dat <- bc %>% select(-id) %>%
  rename(y = diagnosis) %>%
  mutate(y = ifelse(y=="B", 0, 1))

#split <- initial_split(dat, prop = 0.8)

training_df <- dat # split %>% training()

#testing_df <- split %>% testing()
```

Functions

```
loglike_func <- function(dat, betavec){
  # setting up an intercept
  dat_temp = dat %>%
    rename(intercept = y) %>%
    mutate(intercept = rep(1, nrow(dat) ))
  dat_x = unname(as.matrix(dat_temp)) # creating the x matrix

  # finding the pi values
  u = dat_x %*% betavec
  pi <- exp(u) / (1+exp(u))

  # loglikelihood
  loglik <- sum(dat$y*u - log(1 + exp(u)))

  #gradient
```

```

grad <- t(dat_x)%*%(dat$y - pi)

# Hessian
mat_temp = matrix(0, nrow = dim(dat)[1], ncol = dim(dat)[1])
diag(mat_temp)= pi*(1-pi)
hess = -t(dat_x)%*% mat_temp %*%(dat_x)

return(list(loglik = loglik, grad = grad, hess = hess))
}

#loglike_func(dat, betavec = c( rep(0.03, 31))) # test!

NewtonRaphson <- function(dat, start, tol = 1e-11, maxiter = 200){
  i = 0
  cur = start
  stuff = loglike_func(dat, cur)
  res <- c(i=0, "loglik" = stuff$loglik, cur, "step" = 1)
  prevloglik <- -Inf # To make sure it iterates
  while(i < maxiter && abs(stuff$loglik - prevloglik) > tol) {
    step = 1
    i <- i + 1
    prevloglik <- stuff$loglik

    # ascent direct check
    eigen_vals = eigen(stuff$hess)
    if(max(eigen_vals$values) <= 0 ){ # we don't want neg def
      hess = stuff$hess
    } else{ # if it is neg def then need to adjust
      hess = stuff$hess - (max(eigen_vals$values) + 0.1)*diag(nrow(stuff$hess))
    }
    prev <- cur
    cur <- prev - rep(step, length(prev))*(solve(stuff$hess) %*% stuff$grad)
    stuff <- loglike_func(dat, cur) # log-lik, gradient, Hessian

    # doing half stepping
    while(stuff$loglik < prevloglik){
      stuff <- loglike_func(dat, prev)
      step <- step / 2 # this is where half steping happens
      cur <- prev - rep(step, length(prev))*(solve(stuff$hess) %*% stuff$grad)
      stuff <- loglike_func(dat, cur)
    }
    # Add current values to results matrix
    res <- rbind(res, c(i, stuff$loglik, cur, step))
  }

  colnames(res) <- c("intercept", names(dat[, -1]), "i", "loglik", "step")
  return(res)
}

#Running the algorithm with random starting values
betavec = c(rep(0.01, 31))
ans <- NewtonRaphson(training_df, betavec)

```

Results:

##	i	step	loglik
## 1	0.0100000	1.000000e+00	0.010000
## 2	-8.7417254	3.814697e-06	-1.179555
## 3	77.5461663	3.125000e-02	488.328508
## 4	-37.2945192	2.441406e-04	-109.021004
## 5	-0.5525171	3.125000e-02	110.172111
## 6	7.9084537	1.250000e-01	-73.777844
## 7	0.6217574	2.500000e-01	43.286394
## 8	2.8144108	1.000000e+00	49.370660
## 9	6.1210033	1.000000e+00	69.073972
## 10	9.3973114	1.000000e+00	108.752603
## 11	12.7591621	1.000000e+00	165.900414
## 12	18.5983341	1.000000e+00	232.811890
## 13	28.2379507	1.000000e+00	315.879945
## 14	40.9114174	1.000000e+00	410.630445
## 15	61.8581625	1.000000e+00	529.187768
## 16	131.6888689	1.000000e+00	757.177885
## 17	147.4077605	1.250000e-01	807.047155
## 18	149.4996063	1.562500e-02	813.889267
## 19	150.5536412	7.812500e-03	817.351665
## 20	151.0826931	3.906250e-03	819.093312
## 21	151.3477290	1.953125e-03	819.966761
## 22	151.4803745	9.765625e-04	820.404145
## 23	151.5467293	4.882812e-04	820.623002
## 24	151.5633219	1.220703e-04	820.677736
## 25	151.5716188	6.103516e-05	820.705106
## 26	151.5757673	3.051758e-05	820.718792
## 27	151.5758970	9.536743e-07	820.719220
## 28	151.5759132	1.192093e-07	820.719273
## 29	151.5759172	2.980232e-08	820.719286
## 30	151.5759193	1.490116e-08	820.719293
## 31	151.5759203	7.450581e-09	820.719297
## 32	151.5759205	1.862645e-09	820.719297
## 33	151.5759206	4.656613e-10	820.719298
## 34	151.5759206	2.328306e-10	820.719298
## 35	151.5759206	1.164153e-10	820.719298
## 36	151.5759206	5.820766e-11	820.719298
## 37	151.5759206	2.910383e-11	820.719298
## 38	151.5759206	1.455192e-11	820.719298
## 39	151.5759206	7.275958e-12	820.719298
## 40	151.5759206	3.637979e-12	820.719298
## 41	151.5759206	1.818989e-12	820.719298

Variables Name	Beta Coef
intercept	40.000
radius_mean	-11.183
texture_mean	-30.139
perimeter_mean	-44.734
area_mean	0.553
smoothness_mean	0.618
compactness_mean	0.378
concavity_mean	607.372
concave.points_mean	-428.700
symmetry_mean	170.327
fractal_dimension_mean	307.008
radius_se	-154.434
texture_se	254.243
perimeter_se	-15.974
area_se	-5.739
smoothness_se	-6.772
compactness_se	1.201
concavity_se	-110.513
concave.points_se	810.326
symmetry_se	-655.396
fractal_dimension_se	3058.293
radius_worst	-716.255
texture_worst	-8714.163
perimeter_worst	16.337
area_worst	1.128
smoothness_worst	0.452
compactness_worst	-0.118
concavity_worst	-206.369
concave.points_worst	-70.968
symmetry_worst	81.319
fractal_dimension_worst	-50.233

Comparing to glm model

Model Implementation

```
glm_fit <- glm(y~., data=training_df, family = "binomial")
result_glm <- summary(glm_fit)

glm_fit %>% broom::tidy() %>%
  select(term, estimate) %>%
  mutate(estimate = round(estimate, 3)) %>%
  knitr::kable()
```

term	estimate
(Intercept)	-2.881303e+06
radius_mean	2.427012e+06
texture_mean	1.957831e+05
perimeter_mean	1.473188e+06
area_mean	-1.301181e+05
smoothness_mean	-1.524522e+08
compactness_mean	-6.428386e+06
concavity_mean	1.041553e+06
concave.points_mean	-1.715686e+07
symmetry_mean	4.048592e+07
fractal_dimension_mean	-4.232918e+07
radius_se	3.328481e+07
texture_se	6.368392e+06
perimeter_se	1.700716e+06
area_se	-6.393464e+05
smoothness_se	7.491721e+08
compactness_se	-1.773076e+08
concavity_se	1.528648e+08
concave.points_se	-1.259854e+09
symmetry_se	2.890109e+08
fractal_dimension_se	1.512103e+09
radius_worst	-6.130231e+06
texture_worst	-5.832457e+05
perimeter_worst	-3.538204e+05
area_worst	8.950446e+04
smoothness_worst	-2.161128e+07
compactness_worst	8.986336e+06
concavity_worst	-3.027927e+07
concave.points_worst	1.431304e+08
symmetry_worst	-2.473590e+07
fractal_dimension_worst	-3.698324e+07